Ceng790 Big Data Analytics
Assignment 1

Egemen Berk Galatali
2099018@ceng.metu.edu.tr

March 15, 2019

Part I

1. Using the Spark SQL API (accessible with spark.sql("...")), select fields containing the identifier, GPS coordinates, and type of license of each picture.

```
val sqlContext = neworg.apache.spark.sql.SQLContext(spark.sparkContext)
import sqlContext.implicits._

val first = sqlContext.sql("SELECT photo_id, longitude,latitude, license FROM
    flickrMeta")

// For printing DataFrame uncomment below
//first.map(f => "ID: " + f(0) + " - " + "Coord: " + f(1) + "- " + f(2) + " - " +
    "License: " +f(3)).collect().foreach(println)
```

2. Create a DataFrame containing only data of interesting pictures, i.e. pictures for which the license information is not null, and GPS coordinates are valid (not -1.0).

```
val pictures = originalFlickrMeta
    .filter(l => l(15) !=null) // license
    .filter(l => (l.getFloat(10) != 1.0f)) // Coord
    .filter(l =>l.getFloat(11) != -1.0f) // Coord

```

3. Display the execution plan used by Spark to compute the content of this DataFrame(explain()).

```
== Physical Plan ==
*(1) Filter org.apache.spark.sql.catalyst.optimizer.
    CombineTypedFilters$$$Lambda$958/903794242@1e3f0aea.apply
+- *(1) FileScan csv [photo_id#0L,user_id#1,user_nickname#2,date_taken#3,
    date_uploaded#4,device#5,title#6,description#7,user_tags#8,machine_tags#9,
    longitude#10,latitude#11,accuracy#12,url#13,download_url#14,license#15,
    license_url#16,server_id#17,farm_id#18,secret#19,secret_original#20,
    extension_original#21,marker#22] Batched: false, Format: CSV, Location:
    InMemoryFileIndex[file:/home/egemen/IdeaProjects/untitled/flickrSample.txt],
    PartitionFilters: [], PushedFilters: [], ReadSchema: struct<photo_id:bigint,
    user_id:string,user_nickname:string,date_taken:string,date_uploaded:string...

```

4. Display the data of this pictures (show()). Keep in mind that Spark uses lazy execution, so as long as we do not perform any action, the transformations are not executed

photo_id	user_id	user_nickname	date_taken	date_uploaded	device	title	description	user_tags	machine_tags	longi
2860980452	23516515@N07	50mm-traveller	2008-06-19 12:25:...	1221513976	FUJIFILM+FinePix+F30	leaves+in+motion	leaves+in+motion	gallery2flickr	null	-0.0
2445790010	26102073@N08	elibhetluna	2008-04-27 06:27:...	1209302831	null	Animaci%C3%B3n3fg...	null	de,elibhetluna,fotos	null	-0.0
4913556997	34619038@N00	origamidon	2010-08-21 16:24:...	1282422294	null	iPhone.home+%E2%8...	%E2%88%9E+See+the...	accidental,accide...	null	-0.0
1345730799	12402214@N00	Gwyneth+Llewelyn	2006-10-27 13:20:...	1189254028	null	Overview+of+Scope...	Actually%2C+this+...	cleaver,scope,sec...	null	-0.0
1345733105	12402214@N00	Gwyneth+Llewelyn	2006-10-27 13:21:...	1189254077	null	Side+office+of+Sc...	null	cleaver,scope,sec...	null	-0.0
5052929796	9587391@N08	andresoutoarte	2010-10-04 17:30:...	1286238647	null	Lacraia+azul+indigo	Conferir+Padrao+c...	null	null	-0.0
3116901547	50751757@N00	Rojer	2008-12-17 08:07:...	1229578656	EASTMAN+KODAK+COM...	100_9604.JPG	Dove+Court%2C+Ora...	ca,california,chr...	null	-1.
3117729084	50751757@N00	Rojer	2008-12-17 08:07:...	1229578696	EASTMAN+KODAK+COM...	100_9605.JPG	Dove+Court%2C+Ora...	ca,california,chr...	null	-1.
3737526549	46267632@N00	andreweland	2009-07-18 23:13:...	1248069477	NIKON+CORPORATION...	Mount+Diablo+Summit	null	null	null	-0.3
3117764790	50751757@N00	Rojer	2008-12-17 08:17:...	1229580043	EASTMAN+KODAK+COM...	100_9644.JPG	Gingerbread+Town+...	ca,california,chr...	null	-1.
3117768410	50751757@N00	Rojer	2008-12-17 08:18:...	1229580176	EASTMAN+KODAK+COM...	100_9647.JPG	Frosty+Frostys+an...	ca,california,chr...	null	-1.
4262750956	46267632@N00	andreweland	2010-01-05 15:03:...	1263135857	NIKON+CORPORATION...	Tagines%2C+Fez	null	null	null	-0.0
3397220196	35795022@N08	miperricazoe	2009-03-30 02:47:...	1238374073	null	38+im%C3%A1genes+...	002_1.JPG%0ACSC_0...	favoritos	null	-0.0
3117773794	50751757@N00	Rojer	2008-12-17 08:19:...	1229580373	EASTMAN+KODAK+COM...	100_9651.JPG+Ging...	Gingerbread+Town%...	ca,california,chr...	null	-1.
3117761408	50751757@N00	Rojer	2008-12-17 08:16:...	1229579905	EASTMAN+KODAK+COM...	100_9641.JPG	Snow+Carolers.+%0...	ca,california,chr...	null	-1.
4591167499	84031328@N00	Camus+Live+Art	2010-04-14 13:56:...	1273406130	Canon+EOS+7D	Ghana+Tour+2010	Posters+at+the+Wi...	aids,art+educatio...	null	-0.9
4591166029	84031328@N00	Camus+Live+Art	2010-04-15 15:50:...	1273406083	Canon+EOS+7D	Ghana+Tour+2010	War+Dance	aids,art+educatio...	null	-0.8
3765897146	39768211@N07	Chef+Cooke	2009-07-20 11:53:...	1248790291	Canon+EOS+REBEL+T1i	20090720_IcourMil...	Local+restaurant+...	africa,ghana,idds...	null	-0.6
3755727437	39768211@N07	Chef+Cooke	2009-07-20 20:55:...	1248565177	Canon+EOS+REBEL+T1i	20090720_BurkinaF...	Kerosene+lanterns...	africa,ghana,idds...	null	-0.5
8491558947	22898994@N00	svengaarn	2013-02-17 13:10:...	1361374034	Apple+iPhone+4S	Bolgatanga%2C+Ghana	null	dhf,ghana,gspd	null	-0.9

only showing top 20 rows

5. Our goal is now to select the pictures whose license is NonDerivative. To this end we will use a second file containing the properties of each license. Load this file in a DataFrame and do a join operation to identify pictures that are both interesting and NonDerivative. Examine the execution plan and display the results.

We first load the file and then filter this license frame so that we get only the nonDerivative Licenses as a tuple containing name of the license as first parameter and dumy integer as a second parameter so that join operation will perform. We rename the column of the first frame of nonDervLicense so that in the join procedure we ease our job. In the join procedure, we join interesting pictures, which were generated before, with nonDerivative Licenses.

```
val flickrLicense = spark.sqlContext
      .read
      .format("csv")
      .option("delimeter", " ")
      .option("sep", "\t")
      .option("inferSchema", "true")
      .option("header", "true")
      .load("/home/egemen/IdeaProjects/untitled/src/main/scla/edu/metu/ceng790/Assignment1/FlickrLicense.txt")

val nonDervLicense = flickrLicense
                        .filter(lis =>(lis.getInt(3) == 1))
                        .map(l => (l.getString(0), 1))

// We rename the column in order to use Seq("license") so that we remove duplicate columns.
val nonDervs = nonDervLicense.withColumnRenamed("_1","license")

println(flickrLicense.explain()) // For examining the execution
pictures.join(nonDervs, Seq("license")).show() // For showing the results.

```

6. During a work session, it is likely that we reuse multiple time the DataFrame of interesting pictures. It would be a good idea to cache it to avoid recomputing it fromthe file each time we use it. Do this, and examine the execution plan of the join operation again. What do you notice?

When I compared the cached execution plan(Physical Plan) with non-cached version, these two new lines were added to the cached version at the top of the plan.

```
InMemoryTableScan [license#15, photo_id#0L, user_id#1, user_nickname#2,
    date_taken#3, date_uploaded#4 ...]
    +- InMemoryRelation [license#15, photo_id#0L, user_id#1, user_nickname#2,
    date_taken#3, date_uploaded#4 ...], StorageLevel(disk, memory, deserialized,
    1 replicas)

```

7. Save the final result in a csv file (write). Dont forget to add a header to reuse it more easily.

```
val result = pictures.join(nonDervs, Seq("license"))

result.write.format("csv")
        .option("delimeter", " ")
        .option("sep", "\t")
        .option("header", true)
        .save("nonDervLicensedPictures.csv")

```

Results are attached to the submission folder.

Part II

1. Display the 5 lines of the RDD (take(5)) and display the number of elements in the RDD (count())

```
originalFlickrMeta.take(5).foreach(println)
println("Count of the RDD: " + originalFlickrMeta.count())
```

2. Transform the RDD[String] in RDD[Picture] using the Picture class. Only keep interesting pictures having a valid country and tags. To check your program, display 5 elements

```
// We first transform the RDD[String] to RDD[Pictures] using map
// Then filter out the picitures with valid country and tags
val pictures = originalFlickrMeta
                  .map(f => new Picture(f.split("\t")))
                  .filter(f => f.hasValidCountry)
                  .filter(f => f.hasTags)
                  // for printing results uncomment below
                  //.take(5)
                  //.foreach(println)
```

Results are below

```
19/03/13 10:46:45 INFO DAGScheduler: Job 0 finished: take at Part2.scala:37, took 1.129250 s
(UV, aids, art education, ghana, hiv, hiv/aids, hiv prevention, lotos collective, malina de carlo, roberto sanchez-camus, youth visions)
(UV, aids, art education, ghana, hiv, hiv/aids, hiv prevention, lotos collective, malina de carlo, roberto sanchez-camus, youth visions)
(BN, africa, ghana, idds, navrongo)
(UV, africa, ghana, idds, night)
(UV, dhf, ghana, gspd)
```

3. Now group these images by country (groupBy). Print the list of images corresponding to the first country. What is the type of this RDD?

```
pictures.groupBy(f => (f.c, f.toString))
        .take(1)
        .foreach(println)
```

It is actually a tuple contaning the first element as a Key(country) and the second element of it is the String

4. We now wish to process an RDD containing pairs in which the first element is a country, and the second element is the list of tags used on pictures taken in this country. When a tag is used on multiple pictures, it should appear multiple times in the list. As each image has its own list of tags, we need to concatenate these lists, and the flatten function could be useful.

```
val flatted = pictures.map(f => (f.c.toString(), f.userTags))
                      .groupByKey()
                      .map(f => (f._1, f._2.flatten))
                      //.foreach(println)
```

   Mapping operation creates a country name and picture's tags for each picture. After that, We group according to the country name and then flatten the list of tags so that each pictures' tags will be equally distributed to the list.

5. We wish to avoid repetitions in the list of tags, and would rather like to have each tag associated to its frequency. Hence, we want to build a RDD of type RDD[(Country, Map[String, Int])]. The groupBy(identity) function, equivalent to groupBy(x=¿x) could be useful.

```
flatted.map(f => (f._1, f._2.groupBy(x=>x)
                            .map(y => (y._1, y._2.size)))).foreach(println)
```

   Results are below

```
(UV,Map(burkina_faso -> 2, patenschaft -> 2, img_8602.jpg -> 1, community -> 1, zai -> 1, drylands ->
(ML,Map(sand -> 1, canary wharf -> 4, dune -> 1, mezquitas -> 9, tuaregs -> 1, gao -> 2, nomad -> 1, t
(BN,Map(lab -> 5, ghana -> 7, rice -> 1, single mothers -> 1, africa -> 2, idds -> 2, navrongo -> 1))
(AG,Map(3 <- امازيغية ثقافة ,3 <- الطوارق, tamanrasset -> 2, 3 <- الهقار ,3 <- تمنراست, algeria -> 3,
```

6. There are often several ways to obtain a result. The method we used to compute the frequency of tags in each country quickly reaches a state in which the size of the RDD is the number of countries. This can limit the parallelism of the execution as the number of countries is often quite small. Can you propose another way to reach the same result without reducing the size of the RDD until the very end?

   We can fasten the process by calculating the frequency of the tags for each picture before we gather the pictures according to their country.